

Stats101c_hw5

Takao

2022-11-15

Takao Oba

Stats 101C HW 5

Q1) In this Question, we will predict the amount of money expend in college using the other variables in the College data set. Split the data set into a 70% training set and 30% testing set. Use `set.seed(1128)`

```
temp <- read.csv("/Users/takaooba/Downloads/College Fall 2021.csv")
college <- temp[, -c(1,2)]
head(college)
```

##	Private	Apps	Accept	Enroll	Top10perc	Top25perc	F.Undergrad	P.Undergrad
## 1	Yes	1758	1485	419	27	58	2041	174
## 2	No	14463	6166	1757	60	94	8544	671
## 3	Yes	838	651	159	11	25	654	162
## 4	Yes	1127	884	308	30	64	1310	766
## 5	Yes	735	423	366	20	48	2448	707
## 6	Yes	504	482	185	10	36	550	84

##	Outstate	Room.Board	Books	Personal	PhD	Terminal	S.F.Ratio	perc.alumni
## 1	12040	4100	600	1100	92	96	13.2	17
## 2	6550	4598	700	1000	83	100	18.0	15
## 3	8640	3700	400	1915	62	62	12.2	13
## 4	11718	7398	450	1800	73	87	16.4	33
## 5	9210	3782	700	1000	49	51	39.8	15
## 6	9130	3322	450	1450	46	51	12.6	25

##	Grad.Rate	Expend
## 1	72	9060
## 2	80	8055
## 3	48	7634
## 4	76	8871
## 5	34	6562
## 6	54	8686

```
dim(college)
```

```
## [1] 2000 18
```

```
# Splitting 70% training and 30% testing
set.seed(1128)
test.i <- sample(1:2000, 600, replace = F)
college.test <- college[test.i,]

college.train <- college[-test.i,]
```

a) Fit a full multiple linear model using least squares on the training set, and report the MSE obtained using both data sets (training and testing).

```
college.model <- lm(Expend ~ ., data = college.train)
summary(college.model)
```

```
##
## Call:
## lm(formula = Expend ~ ., data = college.train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9501.7 -1401.5  -330.9   854.2 29479.8
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 5690.96704   815.75143    6.976 4.68e-12 ***
## PrivateYes  -509.84234   294.92523   -1.729 0.084083 .
## Apps         0.81070    0.07684   10.551 < 2e-16 ***
## Accept      -1.08476    0.14955   -7.253 6.75e-13 ***
## Enroll       0.96925    0.41746    2.322 0.020390 *
## Top10perc   118.71611   11.03958   10.754 < 2e-16 ***
## Top25perc   -61.16818    8.70870   -7.024 3.38e-12 ***
## F.Undergrad -0.10797    0.07397   -1.460 0.144575
## P.Undergrad -0.06055    0.07456   -0.812 0.416932
## Outstate    0.55630    0.03768   14.765 < 2e-16 ***
## Room.Board  0.02744    0.10672    0.257 0.797134
## Books       1.26521    0.48786    2.593 0.009604 **
## Personal    0.23805    0.13346    1.784 0.074700 .
## PhD        -4.16541    9.09178   -0.458 0.646916
## Terminal    33.12020   10.10478    3.278 0.001073 **
## S.F.Ratio  -291.15598   24.29938  -11.982 < 2e-16 ***
## perc.alumni 12.06622    8.48524    1.422 0.155245
## Grad.Rate  -22.24596    6.23650   -3.567 0.000373 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2875 on 1382 degrees of freedom
## Multiple R-squared:  0.6972, Adjusted R-squared:  0.6934
## F-statistic: 187.2 on 17 and 1382 DF,  p-value: < 2.2e-16
```

The MSE obtained using both data sets

```
pred.train <- predict(college.model, college.train)
pred.test  <- predict(college.model, college.test)

# MSE for the training data set
sum((college.train$Expend - pred.train)^2)/length(pred.train)

## [1] 8160418
```

```
# MSE for the testing data set
sum((college.test$Expend - pred.test)^2)/length(pred.test)

## [1] 10174615
```

b) Fit a ridge regression model on the training set, with lambda chosen by cross-validation. Report the MSE obtained using both data sets (training and testing).

```
set.seed(1128)
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.1.2

## Loading required package: Matrix

## Loaded glmnet 4.1-4
```

```
# Training data
x = model.matrix(Expend ~ ., data = college.train)
x.test <- model.matrix(Expend ~ ., data = college.test)
y = college.train$Expend

i = seq(10, -2, length = 100)
lambda.v = 10^i

model.ridge <- cv.glmnet(x, y, alpha = 0, lambda = lambda.v, type.measure = 'mse', keep=TRUE)
lambda.ridge = model.ridge$lambda.min
lambda.id <- which(model.ridge$lambda == model.ridge$lambda.min)
model.ridge
```

```
##
## Call:  cv.glmnet(x = x, y = y, lambda = lambda.v, type.measure = "mse",      keep = TRUE, alpha = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min   14.2    74 8490594 1358392      17
## 1se 1629.8    57 9696330 1682395      17
```

```
summary(model.ridge)
```

```
##           Length Class  Mode
## lambda           100 -none- numeric
## cvm              100 -none- numeric
## cvsd             100 -none- numeric
## cvup             100 -none- numeric
## cvlo             100 -none- numeric
## nzero            100 -none- numeric
## call              7 -none- call
## name              1 -none- character
## glmnet.fit        12 elnet  list
## fit.preval 140000 -none- numeric
## foldid            1400 -none- numeric
## lambda.min         1 -none- numeric
## lambda.1se         1 -none- numeric
## index             2 -none- numeric
```

```
mod1 <- glmnet(x,y, alpha = 0, lambda = lambda.v)
mod1.train <- predict(mod1, s = lambda.ridge, newx = x)
mod1.test <- predict(mod1, s = lambda.ridge, newx = x.test)
```

```
# Training
mean((mod1.train - college.train$Expend)^2)
```

```
## [1] 8163692
```

```
# Testing
mean((mod1.test - college.test$Expend)^2)
```

```
## [1] 10157872
```

```
# mse <- function(x,y){mean((x-y)^2)}
# mse.1 <- mse(model.ridge$fit[, lambda.id], y)
# mse.1
```

Training MSE is 8163692 and testing MSE is 8163692.

```
# set.seed(1128)
# # Testing
# x = model.matrix(Expend ~ ., data = college.test)
# y = college.test$Expend
#
# i = seq(10, -2, length = 100)
# lambda.v = 10^i
#
# model.ridge <- cv.glmnet(x, y, alpha = 0, lambda = lambda.v, type.measure = 'mse', keep=TRUE)
# lambda.lasso = model.ridge$lambda.min
# lambda.id <- which(model.ridge$lambda == model.ridge$lambda.min)
# model.ridge
# summary(model.ridge)
```

```
#
# mse <- function(x,y){mean((x-y)^2)}
# mse.1 <- mse(model.ridge$fit[, lambda.id], y)
# mse.1
```

c) Fit a lasso model on the training set, with lambda chosen by cross validation. Report the MSE obtained using both data sets (training and testing), along with the number of non-zero coefficient estimates.

```
set.seed(1128)
# Training data
x = model.matrix(Expend ~ ., data = college.train)
x.test = model.matrix(Expend ~ ., data = college.test)
y = college.train$Expend

i = seq(10, -2, length = 100)
lambda.v = 10^i

model.lasso <- cv.glmnet(x, y, alpha = 1, lambda = lambda.v, type.measure = 'mse', keep=TRUE)
lambda.lasso = model.lasso$lambda.min
lambda.id <- which(model.lasso$lambda == model.lasso$lambda.min)
model.lasso
```

```
##
## Call:  cv.glmnet(x = x, y = y, lambda = lambda.v, type.measure = "mse",      keep = TRUE, alpha = 1)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min      0.5      86 8494226 1352101      17
## 1se    305.4      63 9714282 1637899       7
```

```
summary(model.lasso)
```

```
##           Length Class  Mode
## lambda           100 -none- numeric
## cvm              100 -none- numeric
## cvsd             100 -none- numeric
## cvup             100 -none- numeric
## cvlo             100 -none- numeric
## nzero            100 -none- numeric
## call              7 -none- call
## name              1 -none- character
## glmnet.fit        12 elnet  list
## fit.preval    140000 -none- numeric
## foldid          1400 -none- numeric
## lambda.min         1 -none- numeric
## lambda.1se         1 -none- numeric
## index             2 -none- numeric
```

```
mod1 <- glmnet(x,y, alpha = 1, lambda = lambda.v)
mod1.train <- predict(mod1, s = lambda.lasso, newx = x)
mod1.test <- predict(mod1, s = lambda.lasso, newx = x.test)
```

```
# Training
mean((mod1.train - college.train$Expend)^2)
```

```
## [1] 8160593
```

```
# Testing
mean((mod1.test - college.test$Expend)^2)
```

```
## [1] 10170911
```

```
# mse <- function(x,y){mean((x-y)^2)}
# mse.1 <- mse(model.lasso$fit[, lambda.id], y)
# mse.1
```

Training MSE is 8160593 and testing MSE is 10170911.

```
# set.seed(1128)
# # Testing
# x = model.matrix(Expend ~ ., data = college.test)
# y = college.test$Expend
#
# i = seq(10, -2, length = 100)
# lambda.v = 10^i
#
# model.ridge <- cv.glmnet(x, y, alpha = 1, lambda = lambda.v, type.measure = 'mse', keep=TRUE)
# lambda.lasso = model.ridge$lambda.min
# lambda.id <- which(model.ridge$lambda == model.ridge$lambda.min)
# model.ridge
# summary(model.ridge)
#
#
#
# mse <- function(x,y){mean((x-y)^2)}
# mse.1 <- mse(model.ridge$fit[, lambda.id], y)
# mse.1
```

Q2) Use the same data sets in Question 1 to:

a) Fit a PCR model on the training set, with M principal components chosen by cross validation. Report the MSE obtained using both data sets (training and testing). along with the value of M principal components selected by cross-validation. Report the amount of variation explained in the X matrix by those M principal component.


```
# install.packages("pls")
library(pls)
```

```
## Warning: package 'pls' was built under R version 4.1.2
```

```
##
## Attaching package: 'pls'
```

```
## The following object is masked from 'package:stats':
##
## loadings
```

```
pcr.college <- pcr(Expend ~., data = college.train, scale = TRUE, validation = "CV")
validationplot(pcr.college, val.type = "MSEP")
```



From the graph above, we can see that the value of M principal components by cross validation is 2. We will further complete our computations.

Next, we will report the MSE obtained using both data sets (training and testing)

```
# Testing
pcr.pred.test <- predict(pcr.college, college.test, ncomp = 7)
mean((pcr.pred.test - college.test$Expend)^2)
```

```
## [1] 12321549
```

```
# Training
```

```
pcr.pred.train <- predict(pcr.college, college.train, ncomp = 7)
mean((pcr.pred.train - college.train$Expend)^2)
```

```
## [1] 10794544
```

Now that we have assessed the MSEs of both data sets, we will look into the amount of variation explained in the X matrix by those M principal component.


```
summary(pcr.college)
```

```
## Data:      X dimension: 1400 17
## Y dimension: 1400 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              5195    5156    3489    3400    3398    3390    3318
## adjCV           5195    5161    3489    3399    3397    3390    3318
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          3303    3279    3267    3224    3224    3106    3099
## adjCV        3301    3278    3266    3223    3223    3105    3098
##      14 comps 15 comps 16 comps 17 comps
## CV           3026    2930    2910    2891
## adjCV         3029    2928    2909    2889
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          30.870   59.79   66.49   72.28   77.71   82.34   85.77   89.02
## Expend      1.662   55.00   57.29   57.36   57.53   59.47   59.94   60.56
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          91.91   94.31   96.28   97.46   98.40   99.04   99.64
## Expend      60.96   62.03   62.06   64.85   65.02   66.55   68.75
##      16 comps 17 comps
## X           99.87  100.00
## Expend       69.33   69.72
```

Using 85% as our threshold for the amount of variation in the X matrix, we will have that the variation explained is with 7 components and 59.94% of the variation explained.

b) Fit a PLS model on the training set, with M principal components chosen by cross validation. Report the MSE obtained using both data sets (training and testing), along with the value of M principal components selected by cross-validation. Report the amount of variation explained in the X matrix by those M principal component. Note: Use 85% as your threshold for the amount of variation in the X matrix

```
pls.college <- pls(Epend ~ ., data = college.train, scale = TRUE, validation = "CV")  
validationplot(pls.college, val.type = "MSEP")
```



From the graph above, we can see that the value of M principal components by cross validation is 2. We will further complete our computations.

Next, we will report the MSE obtained using both data sets (training and testing)

```
# Testing  
pls.pred.test <- predict(pls.college, college.test, ncomp = 10)  
mean((pls.pred.test - college.test$Epend)^2)
```

```
## [1] 10189596
```

```
# Training
pls.pred.train <- predict(pls.college, college.train, ncomp = 10)
mean((pls.pred.train - college.train$Expend)^2)
```

```
## [1] 8177409
```

Now that we have assessed the MSEs of both data sets, we will look into the amount of variation explained in the X matrix by those M principal component.

```
summary(pls.college)
```

```
## Data:      X dimension: 1400 17
## Y dimension: 1400 1
## Fit method: kernelppls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              5195    3400    3154    3107    3012    2948    2923
## adjCV           5195    3400    3151    3113    3010    2946    2921
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          2910    2908    2905    2901    2900    2900    2899
## adjCV        2908    2906    2903    2900    2898    2898    2897
##      14 comps 15 comps 16 comps 17 comps
## CV          2899    2899    2899    2899
## adjCV        2897    2897    2897    2897
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          28.9    35.75   63.11   68.62   71.83   74.83   77.76   81.22
## Expend      57.5    63.80   64.72   67.28   68.69   69.17   69.39   69.50
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          84.03    86.55   90.02   92.86   94.30   96.19   97.96
## Expend      69.57    69.65   69.70   69.70   69.71   69.71   69.72
##      16 comps 17 comps
## X          99.30   100.00
## Expend      69.72    69.72
```

Using 85% as our threshold for the amount of variation in the X matrix, we will have that the variation explained is with 10 components and 69.65% of the variation explained.

Q3) This question relates to the College data sets in question 1.

(a) Using “Expend” as the response and the other variables as predictors, perform backward stepwise selection (Choose BIC as your criteria) on the training set in order to identify a satisfactory model that uses just a subset of the predictors.

```

library(leaps)


regfit.bck <- regsubsets(Expend ~ ., data = college.train, method = "backward")
summary(regfit.bck)

## Subset selection object
## Call: regsubsets.formula(Expend ~ ., data = college.train, method = "backward")
## 17 Variables (and intercept)
##              Forced in Forced out
## PrivateYes      FALSE      FALSE
## Apps            FALSE      FALSE
## Accept          FALSE      FALSE
## Enroll          FALSE      FALSE
## Top10perc       FALSE      FALSE
## Top25perc       FALSE      FALSE
## F.Undergrad     FALSE      FALSE
## P.Undergrad     FALSE      FALSE
## Outstate        FALSE      FALSE
## Room.Board      FALSE      FALSE
## Books           FALSE      FALSE
## Personal        FALSE      FALSE
## PhD             FALSE      FALSE
## Terminal        FALSE      FALSE
## S.F.Ratio       FALSE      FALSE
## perc.alumni     FALSE      FALSE
## Grad.Rate       FALSE      FALSE
## 1 subsets of each size up to 8
## Selection Algorithm: backward
##              PrivateYes Apps Accept Enroll Top10perc Top25perc F.Undergrad
## 1  ( 1 ) " "          " " " " " " " " " " " "
## 2  ( 1 ) " "          " " " " " " "*" " " "
## 3  ( 1 ) " "          " " " " " " "*" " " "
## 4  ( 1 ) " "          "*" " " " " "*" " " "
## 5  ( 1 ) " "          "*" "*" " " "*" " " "
## 6  ( 1 ) " "          "*" "*" " " "*" "*" " "
## 7  ( 1 ) " "          "*" "*" " " "*" "*" " "
## 8  ( 1 ) " "          "*" "*" " " "*" "*" " "
##              P.Undergrad Outstate Room.Board Books Personal PhD Terminal S.F.Ratio
## 1  ( 1 ) " "          "*" " " " " " " " " " "
## 2  ( 1 ) " "          "*" " " " " " " " " " "
## 3  ( 1 ) " "          "*" " " " " " " " " "*"
## 4  ( 1 ) " "          "*" " " " " " " " " "*"
## 5  ( 1 ) " "          "*" " " " " " " " " "*"
## 6  ( 1 ) " "          "*" " " " " " " " " "*"
## 7  ( 1 ) " "          "*" " " " " " " " "*" "*"
## 8  ( 1 ) " "          "*" " " " " " " " "*" "*"
##              perc.alumni Grad.Rate
## 1  ( 1 ) " "          " "
## 2  ( 1 ) " "          " "
## 3  ( 1 ) " "          " "
## 4  ( 1 ) " "          " "
## 5  ( 1 ) " "          " "
## 6  ( 1 ) " "          " "

```

```
## 7 ( 1 ) " " " "
## 8 ( 1 ) " " "*"
```

```
plot(regfit.bck, scale = "bic")
```



```
# out <- summary(regsubsets(Expend ~ ., data = college.train, method = "backward"))
# qplot(1:10, out$bic) + geom_line()
```

We assess which predictors to examine through looking at the predictors where the bar touches the top of the graph which are Apps, Accept, Top10perc, Top25perc, Outstate, Terminal, S.F.Ratio, Grad.Rate

```
# The satisfactory model is below
college.train.1 <- college.train[,c(2,3,5,6,9,14,15,17,18)]
college.model <- lm(Expend ~ ., data = college.train.1)
summary(college.model)
```

```
##
## Call:
## lm(formula = Expend ~ ., data = college.train.1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10166.6  -1437.4   -258.8    793.6   29474.3
##
```

```
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) 6630.17593   664.30790    9.981 < 2e-16 ***
## Apps         0.82166     0.07418   11.077 < 2e-16 ***
## Accept       -0.91879     0.10855   -8.464 < 2e-16 ***
## Top10perc    127.59278    10.51691   12.132 < 2e-16 ***
## Top25perc   -63.46323     8.66851   -7.321 4.15e-13 ***
## Outstate      0.52331     0.02873   18.216 < 2e-16 ***
## Terminal     33.81667     6.57410    5.144 3.08e-07 ***
## S.F.Ratio   -292.72328    23.63735  -12.384 < 2e-16 ***
## Grad.Rate   -23.29867     5.89937   -3.949 8.23e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2891 on 1391 degrees of freedom
## Multiple R-squared:  0.6919, Adjusted R-squared:  0.6901
## F-statistic: 390.4 on 8 and 1391 DF,  p-value: < 2.2e-16
```

To find the MSE of the backwards stepwise BIC model, we have

```
pred.train <- predict(college.model, college.train)
pred.test  <- predict(college.model, college.test)
```

```
# MSE for the training data set
sum((college.train$Expend - pred.train)^2)/length(pred.train)
```

```
## [1] 8302748
```

```
# MSE for the testing data set
sum((college.test$Expend - pred.test)^2)/length(pred.test)
```

```
## [1] 10200996
```

The MSE for training is 8302748 and the MSE for testing is 10200996

(b) Fit a GAM on the training data, using “Expend” as the response and the features selected in the previous step (Part a) as your predictors. Plot the results, and explain your findings.

```
library(splines)
library(ISLR)
# install.packages("gam")
library(gam)
```

```
## Warning: package 'gam' was built under R version 4.1.2
```

```
## Loading required package: foreach
```

```
## Warning: package 'foreach' was built under R version 4.1.2
```


```
## Loaded gam 1.22
```

```
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 4.1.2
```


```
ggplot(data= college.train.1, aes(Apps, Expend)) + geom_point() + geom_smooth()
```

```
## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```




```
ggplot(data= college.train.1, aes(Accept, Expend)) + geom_point() + geom_smooth()
```

```
## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```




```
ggplot(data= college.train.1, aes(Top10perc, Expend)) + geom_point() + geom_smooth()
```

```
## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```




```
ggplot(data= college.train.1, aes(Top25perc, Expend)) + geom_point() + geom_smooth()
```

```
## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```


```
ggplot(data= college.train.1, aes(Outstate, Expend)) + geom_point() + geom_smooth()
```

```
## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```




```
ggplot(data= college.train.1, aes(Terminal, Expend)) + geom_point() + geom_smooth()
```

```
## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```




```
ggplot(data= college.train.1, aes(S.F.Ratio, Expend)) + geom_point() + geom_smooth()
```

```
## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```



```
ggplot(data= college.train.1, aes(Grad.Rate, Expend)) + geom_point() + geom_smooth()
```

```
## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```



```
gam0 <- lm(Expend ~ bs(Apps, 4) + Accept + Top10perc + bs(Top25perc, 4) + Outstate + bs(Terminal, 4) + bs(S.F.Ratio, 4) + bs(Grad.Rate, 4), data = college.train)
summary(gam0)
```

```
##
## Call:
## lm(formula = Expend ~ bs(Apps, 4) + Accept + Top10perc + bs(Top25perc,
##      4) + Outstate + bs(Terminal, 4) + bs(S.F.Ratio, 4) + bs(Grad.Rate,
##      4), data = college.train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -13447.8  -1126.0   -296.1    756.8  29939.3
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.569e+04  2.185e+03   7.180 1.13e-12 ***
## bs(Apps, 4)1    8.653e+02  5.204e+02   1.663  0.09663 .
## bs(Apps, 4)2    5.133e+03  1.047e+03   4.905 1.05e-06 ***
## bs(Apps, 4)3    8.705e+03  1.747e+03   4.984 7.00e-07 ***
## bs(Apps, 4)4    7.117e+03  2.363e+03   3.012  0.00264 **
## Accept        -4.744e-01  1.102e-01  -4.305 1.79e-05 ***
## Top10perc       6.254e+01  1.160e+01   5.391 8.22e-08 ***
## bs(Top25perc, 4)1 -1.932e+03  1.199e+03  -1.612  0.10728
## bs(Top25perc, 4)2 -1.477e+03  7.993e+02  -1.848  0.06475 .
## bs(Top25perc, 4)3 -4.713e+03  1.144e+03  -4.121 4.00e-05 ***
```

```
## bs(Top25perc, 4)4 -5.378e+02 1.231e+03 -0.437 0.66220
## Outstate 4.985e-01 2.844e-02 17.527 < 2e-16 ***
## bs(Terminal, 4)1 -2.985e+03 1.920e+03 -1.554 0.12037
## bs(Terminal, 4)2 -2.672e+02 1.071e+03 -0.250 0.80295
## bs(Terminal, 4)3 -5.012e+02 1.319e+03 -0.380 0.70408
## bs(Terminal, 4)4 1.203e+03 1.226e+03 0.981 0.32659
## bs(S.F.Ratio, 4)1 -9.980e+03 1.385e+03 -7.206 9.46e-13 ***
## bs(S.F.Ratio, 4)2 -1.716e+04 1.201e+03 -14.289 < 2e-16 ***
## bs(S.F.Ratio, 4)3 -1.486e+04 1.982e+03 -7.498 1.16e-13 ***
## bs(S.F.Ratio, 4)4 -1.462e+04 1.712e+03 -8.539 < 2e-16 ***
## bs(Grad.Rate, 4)1 4.054e+03 1.929e+03 2.102 0.03577 *
## bs(Grad.Rate, 4)2 9.817e+02 1.168e+03 0.840 0.40097
## bs(Grad.Rate, 4)3 1.365e+03 1.780e+03 0.767 0.44343
## bs(Grad.Rate, 4)4 2.576e+03 1.874e+03 1.375 0.16941
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2668 on 1376 degrees of freedom
## Multiple R-squared: 0.7404, Adjusted R-squared: 0.7361
## F-statistic: 170.7 on 23 and 1376 DF, p-value: < 2.2e-16
```

The lowest p-value polynomials are 3 for Apps, 3 for Top25perc, 1 for Terminal, 2 for S.F.Ratio, and 1 for Grad.Rate




(c) Evaluate the model obtained on the testing data set, and explain the results obtained.


We will determine what coefficients to use for from the previous summary output. For the non linear components Apps, Top25perc, Terminal, S.F.Ratio, and Grad.Rate, we will utilize the polynomial with the lowest p-value. Thus, we will have the following:

```
gam1 <- gam(Expend ~ s(Apps,3) + Accept + Top10perc + s(Top25perc,3) + Outstate + s(Terminal,1) + s(S.F.Ratio,2) + s(Grad.Rate,1))
# gam1 <- gam(Expend ~ s(Apps) + s(Accept) + s(Top10perc) + s(Top25perc) + s(Outstate) + s(Terminal) + s(S.F.Ratio) + s(Grad.Rate))

par(mfrow = c(1,3))
plot(gam1, se = TRUE, col = "blue")
```

```
## Warning in pf(nl.chisq/nldf, nldf, rdf): NaNs produced
```





```
pred <- predict(gam1, newdata = college.test)
```


```
# We will evaluate the model obtained on the testing data set by looking at the MSE.
```

```
# The MSE can be determined through below
```

```
MSE <- mean((college.test$Expend - pred)^2)
```

```
MSE
```

```
## [1] 8938132
```

```
# fit <- lm(Expend ~ bs(Apps, knots = c(25,40,60)), data = college)
# dim(bs(college$Expend, knots = c(25,40, 60)))

# fit.1 = lm(Expend~Apps, data = college)
# fit.2 = lm(Expend~poly(Apps,2), data = college)
# fit.3 = lm(Expend~poly(Apps,3), data = college)
# fit.4 = lm(Expend~poly(Apps,4), data = college)
# fit.5 = lm(Expend~poly(Apps,5), data = college)
# anova(fit.1,fit.2,fit.3,fit.4,fit.5)
```

##(d) For which variables, if any, is there evidence of a non-linear relationship with the response?

The variables with a non-linear relationship with the response are Apps, Top25perc, Terminal, S.F.Ratio, Grad.Rate

Q4) Comment on the results obtained. How accurately can we predict the amount of money expend by college students? Is there much difference among the testing MSEs resulting from these seven approaches?


The MSEs for the seven approaches are as follows 1. Least Squares full model using lm 10174615 2. Ridge model with the best lambda 10157872 3. Lasso model with the best lambda 10170911 4. PCR model 12321549 5. PLS model 10189596 6. Stepwise backward regression using BIC 10200996 7. GAM 8938132

```
temp <- c(10174615,10157872,10170911,12321549,10189596,10200996,8938132)

temprow <- c("Least Squares", "Ridge Model", "Lasso Model", "PCR", "PLS", "Backward BIC", "GAM")
tempdf <- data.frame("Models" <- temprow, "MSEs" <- temp)

library(ggplot2)
ggplot(data = tempdf, aes(Models, MSEs)) + geom_bar(stat= "identity", color = "blue", fill = "lightblue")
```

MSEs of each Models



In the table above, we are able to see the various MSEs by the seven models that we have generated. We can see that the GAM model has the lowest MSE, the PCR model has the highest MSE, and the other models have fairly similar MSE values.

Q5) You may have seen the “betterbirths2000” data in Stats 10. It consists of a random sample of 2000 births in North Carolina that are collected in order to track health issues in new born babies. These data are saved in a file better2000births.csv on bruinlearn Week 8.

```
temp <- read.csv("/Users/takaooba/Downloads/better2000births.csv")
births <- na.omit(temp)
```

```

births$Marital <- factor(births$Marital)
births$Racemom <- factor(births$Racemom)
births$Racedad <- factor(births$Racedad)
births$Hispmom <- factor(births$Hispmom)
births$Hispdad <- factor(births$Hispdad)
births$Habit <- factor(births$Habit)
births$MomPriorCond <- factor(births$MomPriorCond)
births$BirthDef <- factor(births$BirthDef)
births$DelivComp <- factor(births$DelivComp)
births$BirthComp <- factor(births$BirthComp)
head(births)

```

```

##   Gender Premie weight Apgar1 Fage Mage Feduc Meduc TotPreg Visits Marital
## 1   Male    No   124     8   31  25   13   14         1    13   Married
## 2 Female    No   177     8   36  26    9   12         2    11 Unmarried
## 3   Male    No   107     3   30  16   12    8         2    10 Unmarried
## 4 Female    No   144     6   33  37   12   14         2    12 Unmarried
## 5   Male    No   117     9   36  33   10   16         2    19   Married
## 6 Female    No    98     4   31  29   14   16         3    20   Married
##   Racemom Racedad Hispmom Hispdad Gained   Habit MomPriorCond BirthDef
## 1   White   White NotHisp NotHisp    40 NonSmoker      None      None
## 2   White   White Mexican Mexican    20 NonSmoker      None      None
## 3   White Unknown Mexican Unknown    70 NonSmoker At Least One      None
## 4   White   White NotHisp NotHisp    50 NonSmoker      None      None
## 5   White   Black NotHisp NotHisp    40 NonSmoker At Least One      None
## 6   White   White NotHisp NotHisp    21 NonSmoker      None      None
##   DelivComp BirthComp
## 1 At Least One      None
## 2 At Least One      None
## 3 At Least One      None
## 4 At Least One      None
## 5           None      None
## 6           None      None

```

a) Split your data into Training and Testing. You should have 1000 observations in your training data after omitting the missing values in your data. Use the set.seed “1128” to do the split. Use a tree (not pruned) to predict whether a baby will be born prematurely or normal. What is the testing misclassification error?

```

dim(births)

## [1] 1998  21

set.seed(1128)
test.i <- sample(1:nrow(births), 1000, replace = F)

births.test <- births[-test.i,]
births.train <- births[test.i,]

```

Using a tree


```
# install.packages("tree")
library(tree)
```

```
## Warning: package 'tree' was built under R version 4.1.2
```

```
births.model <- tree(formula = factor(Premie) ~ ., data = births.train)
```

```
## Warning in tree(formula = factor(Premie) ~ ., data = births.train): NAs
## introduced by coercion
```

```
plot(births.model)
text(births.model, pretty = 0)
```



```
births.y <- births.test$Premie
```

```
preds <- predict(births.model, newdata = births.test, type = "class")
```

```
## Warning in pred1.tree(object, tree.matrix(newdata)): NAs introduced by coercion
```

```
conf.matrx <- table(preds, factor(births.test$Premie))
```

```
conf.matrx
```

```
##  
## preds  No Yes  
##    No  895  49  
##    Yes   13  41
```

```
(conf.matrx[2,1] + conf.matrx[1,2])/sum(conf.matrx)
```

```
## [1] 0.06212425
```

b) Use cross-validation to determine if the tree can be improved through pruning. If so, prune the tree to the appropriate size and provide a plot.

```
cv.train <- cv.tree(births.model, FUN = prune.misclass)
```

```
## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by  
## coercion
```

```
## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion
```

```
## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by  
## coercion
```

```
## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion
```

```
## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by  
## coercion
```

```
## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion
```

```
## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by  
## coercion
```

```
## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion
```

```
## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by  
## coercion
```

```
## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion
```

```
## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by  
## coercion
```

```
## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion
```

```
## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by
## coercion

## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion

## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by
## coercion

## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion

## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by
## coercion

## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion


## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by
## coercion

## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion
```


```
names(cv.train)
```

```
## [1] "size" "dev" "k" "method"
```

```
# Find best value using the plot function
plot(cv.train$dev ~ cv.train$size, type = "b")
```



```
pruned.fit <- prune.misclass(births.model, best = 3)
plot(pruned.fit)
text(pruned.fit, pretty = TRUE)
```



We will see if pruning makes the misclassification rate better. Recall that without pruning, we had a misclassification rate of 0.06212425

```
summary(pruned.fit)
```

```
##
## Classification tree:
## snip.tree(tree = births.model, nodes = c(4L, 5L, 3L))
## Variables actually used in tree construction:
## [1] "weight"
## Number of terminal nodes: 3
## Residual mean deviance: 0.322 = 321.1 / 997
## Misclassification error rate: 0.054 = 54 / 1000
```

```
# births.train
```

Now, we see that the misclassification rate is 0.056 which is lower than the misclassification rate without pruning.

c) Interpret your pruned tree (or your tree in (a) if you did not need to prune). In particular, does it tell us whether smoking is a potential cause of premature births? What factors are associated with premature births?

Basing on the pruned tree, we have that the factors that are associated with premature births is Weights. The pruned plot have performed fairly well with a low misclassification rate. Intuitively, smoking should be a potential cause of premature births, but in this decision tree, we were not able to conclude that smoking is directly associated with premature births. However, according to CDC.gov ([https://www.cdc.gov/tobacco/campaign/tips/diseases/pregnancy.html#:~:text=Smoking%20slows%20your%20baby's%20growth,babies%20often%20have%20a%20higher%20risk%20of%20premature%20births,which%20can%20lead%20to%20low%20birth%20weight%20babies,which%20are%20at%20higher%20risk%20of%20sudden%20infant%20death%20syndrome%20\(SIDS\).](https://www.cdc.gov/tobacco/campaign/tips/diseases/pregnancy.html#:~:text=Smoking%20slows%20your%20baby's%20growth,babies%20often%20have%20a%20higher%20risk%20of%20premature%20births,which%20can%20lead%20to%20low%20birth%20weight%20babies,which%20are%20at%20higher%20risk%20of%20sudden%20infant%20death%20syndrome%20(SIDS).)) we have that smoking slows the baby's growth before birth, or more specifically, the weight of the baby. The factor that is most associated with premature births is Weights.

d) What is the testing misclassification error rate of your pruned tree? Keep in mind that approximately 9% of all births are premature. This means that if a doctor simply predict “not premature” ALWAYS, he or she will have only a 9% misclassification error. Did you do better based on your tree models?

```
preds <- predict(pruned.fit, newdata = births.test, type = "class")

## Warning in pred1.tree(object, tree.matrix(newdata)): NAs introduced by coercion

conf.matrx <- table(preds, factor(births.test$Premie))

conf.matrx

##
## preds   No  Yes
##    No  901  50
##    Yes   7  40

(conf.matrx[2,1] + conf.matrx[1,2])/sum(conf.matrx)

## [1] 0.05711423
```

The misclassification rate for testing is 0.05711423. Since the misclassification rate is 5.71%, which is less than the 9% misclassification error, we have that this tree models is better than simply predicting “not premature” for all the babies.